Neural Net
Final Report

AllFresh
Reporter: o100442 Wei-Yuan Hsu
Advisor: Prof. Hahn-Ming Lee

Outline

Rank

Data

Model

Results

Discussion

Rank

Rank

Highest rank	200
Public score	294
Specialized Prize for the last 10 days	444
Specialized Prize for the second-day prediction	517
Last rank	523

Problem Statement

- Pollutant Prediction
  - PM2.5, PM10, O3 (Beijing only)

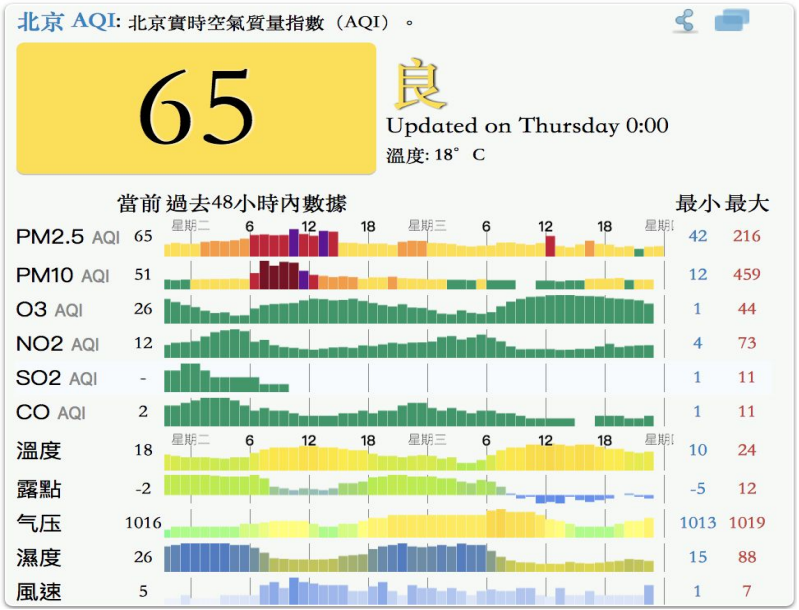

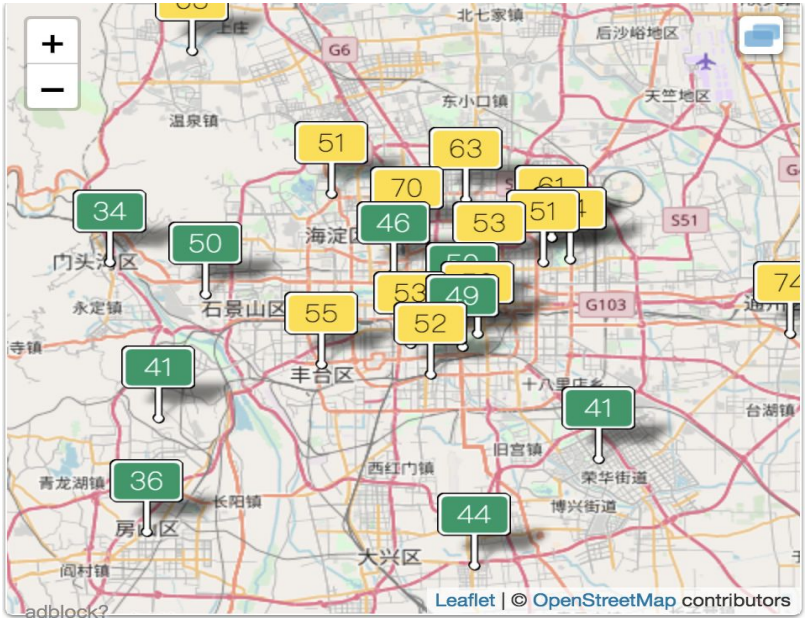

Platform

- GPU Server
- Pandas, Numpy, Keras

Data

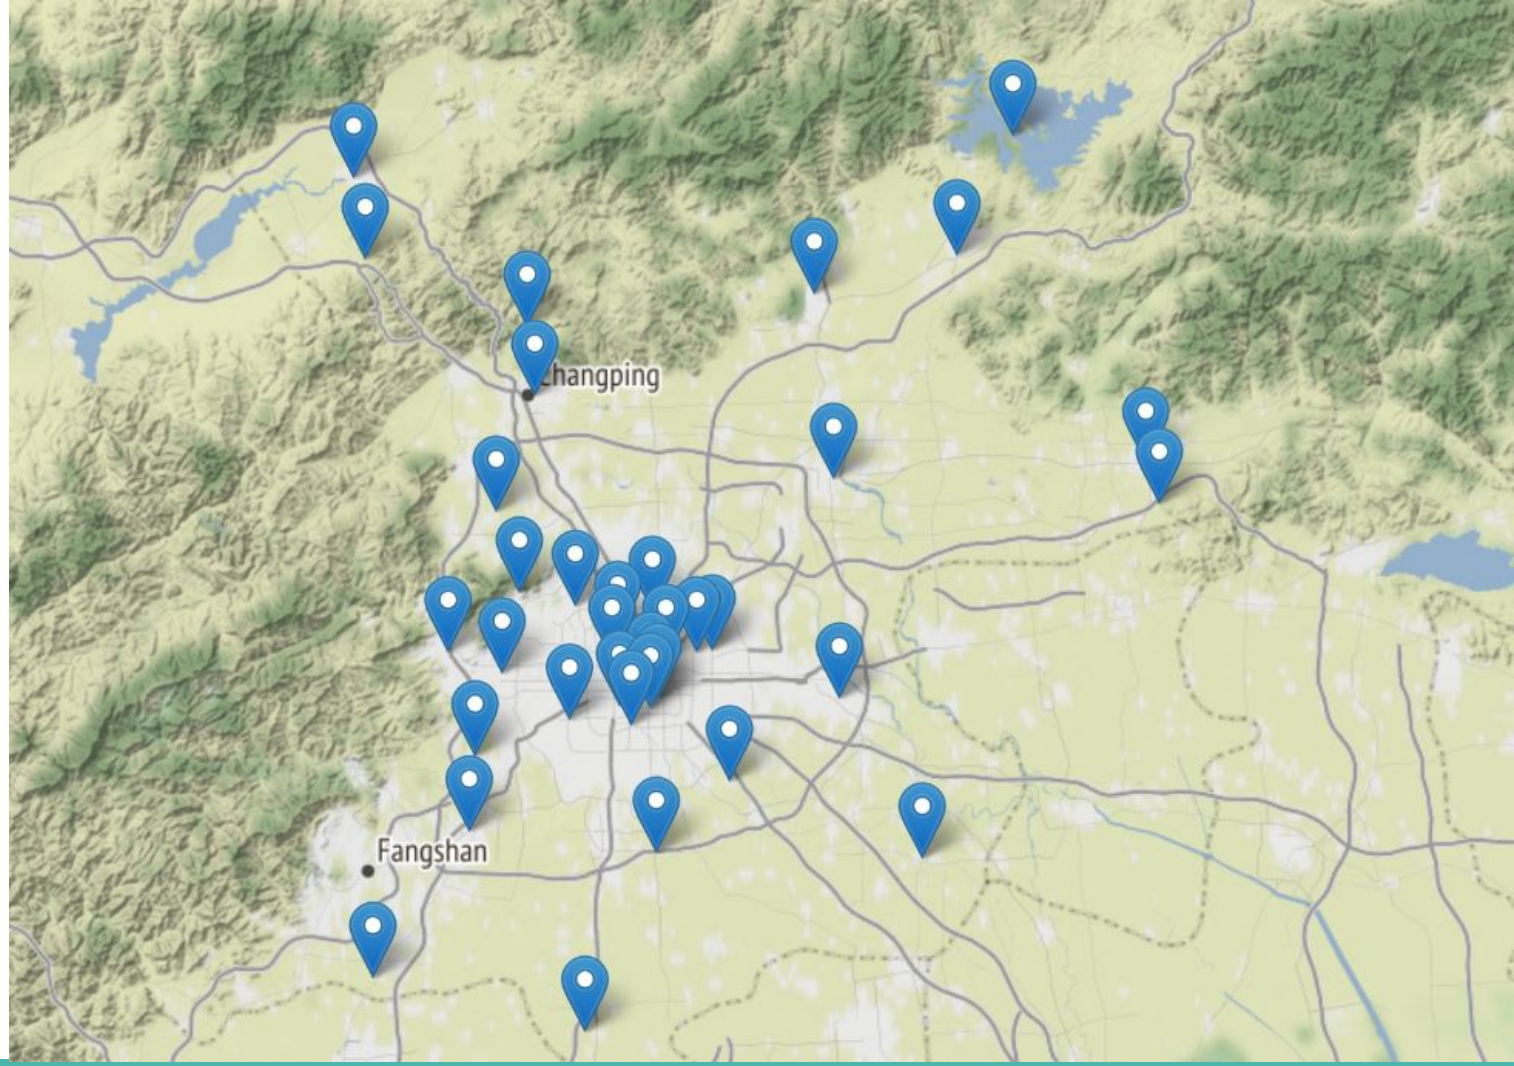

Input Data - Data Preprocessing

- Remove outlier (ex. Temperature, wind direction)
- Denoising: median number

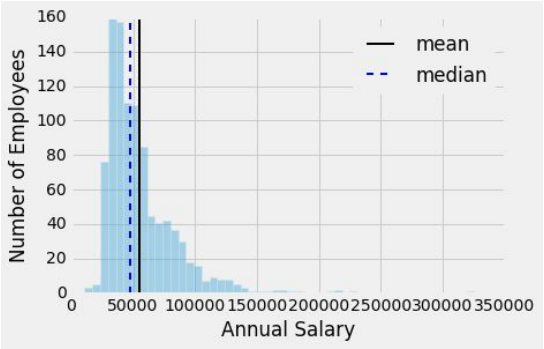

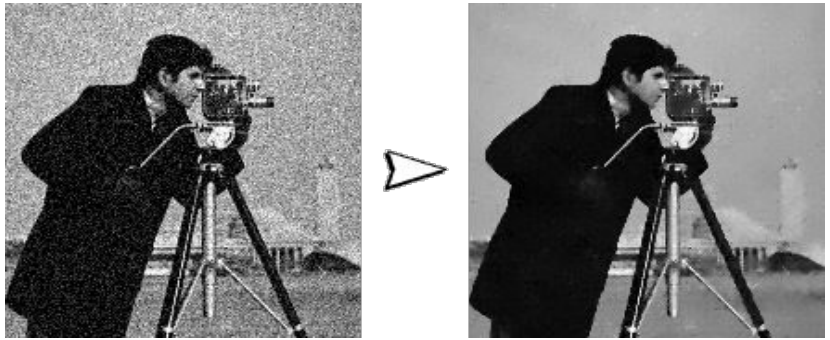

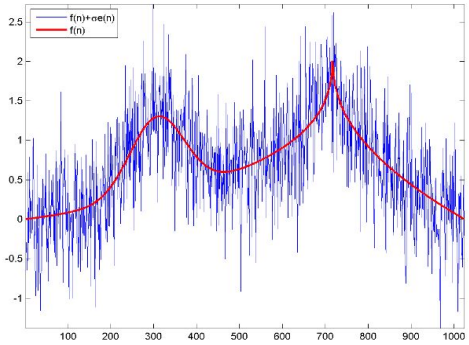

Reference:

1. Local Approximations in Signal and Image Processing, http://www.cs.tut.fi/~lasip/2D/
2. Wavelet Denoising and Nonparametric Function Estimation, https://goo.gl/HJWVrs

Input Data - Handling missing data

- Missing value
  - missing # <= 23 hrs - pandas interpolation
  - missing 1 day (24 hrs) - average of previous 3 days
- External data
  - beijingair

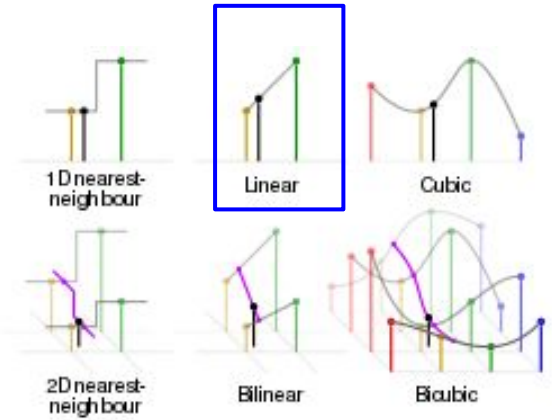

Feature Engineering

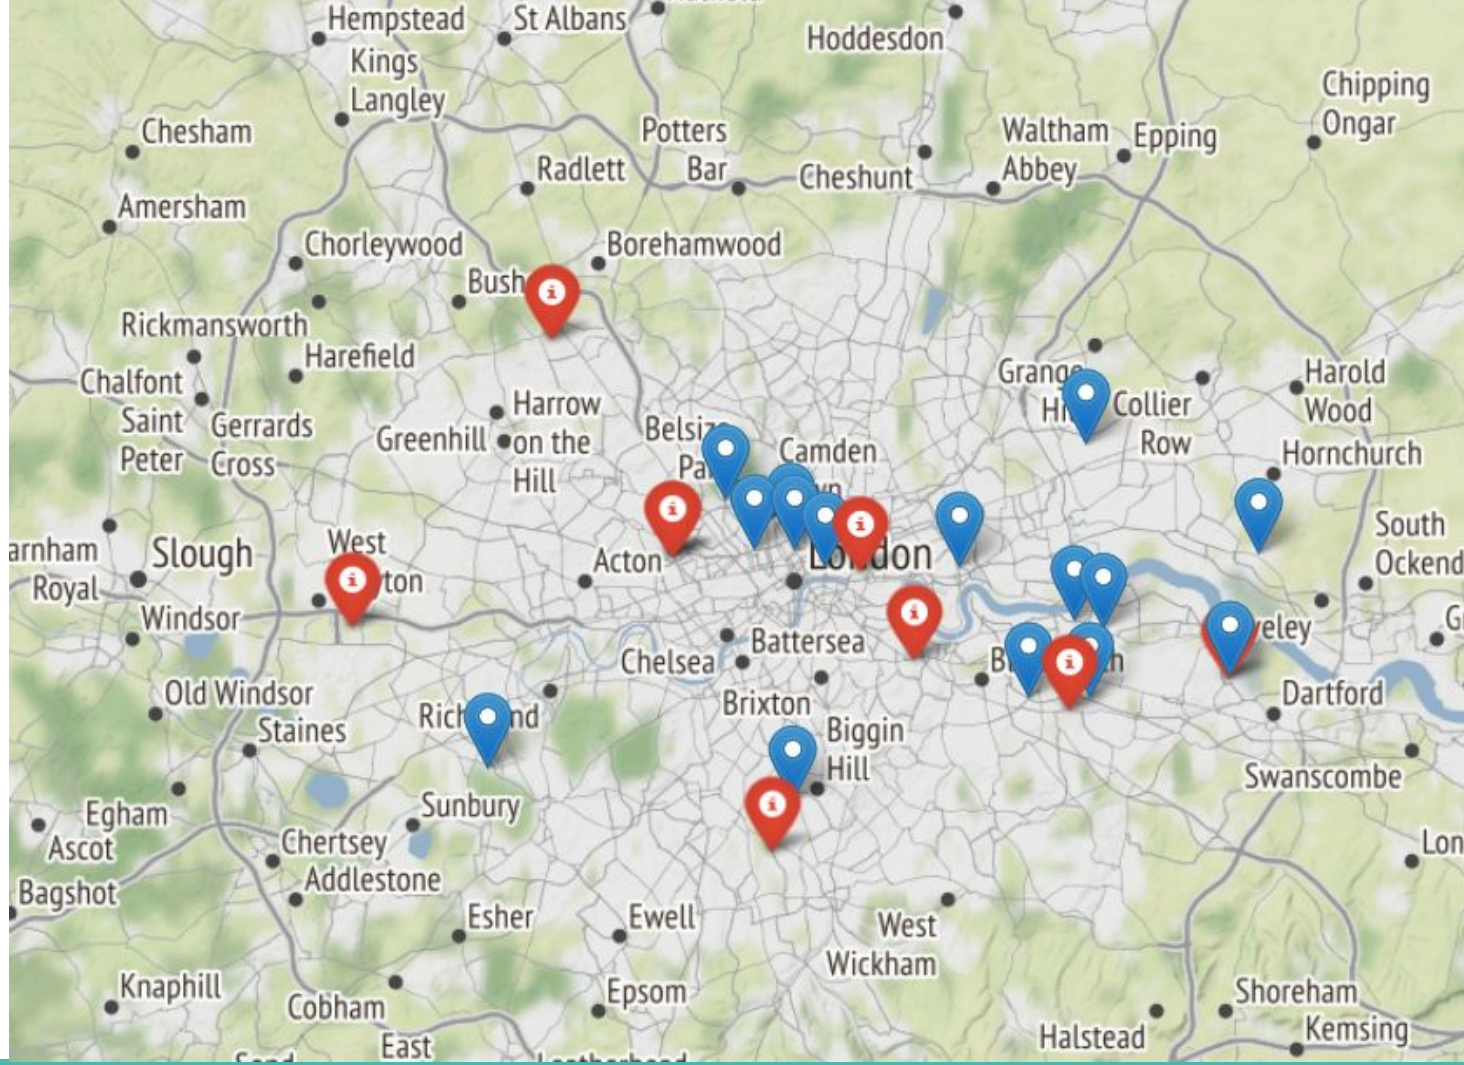

Feature engineering

RNN is powerful enough to learn the pattern behind features.

- **Important feature** (what we want to predict)
  - PM2.5
  - PM10
  - O3
- Other pollution
  - NO2
  - CO
  - SO2
- Weather information
  - temperature
  - pressure
  - humidity
  - wind_direction
  - wind_speed/kph

Reference
[1] Arturus/kaggle-web-traffic

Model

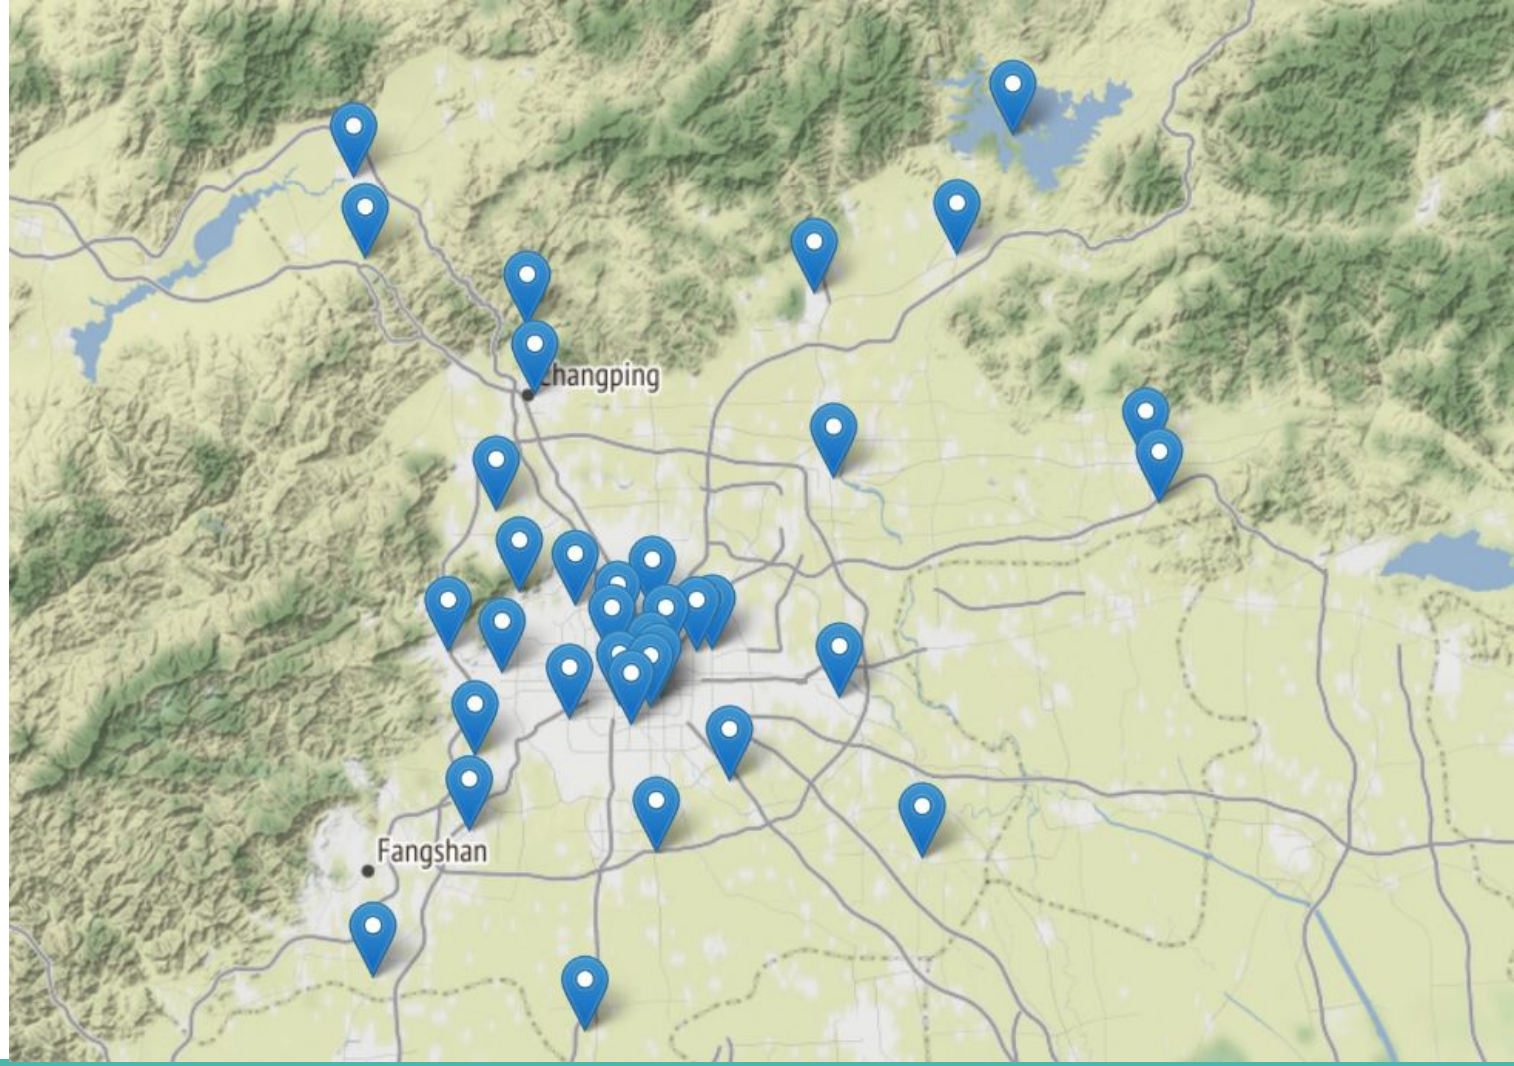

Training and Validation

- SMAPE loss boosting
  - Clipping
- Regularization
  - Overfit -> just fit
  - Neural Network: Dropout (20%)
- Early Stopping
  - Stop the rest training iteration at early stage if the model in case time waste
  - patience: 6

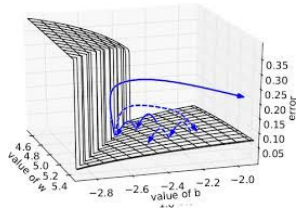

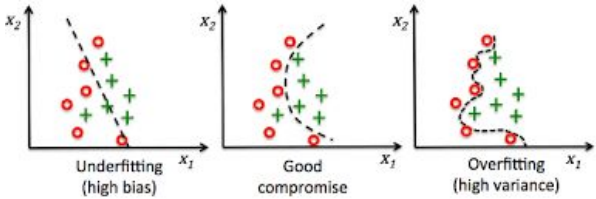

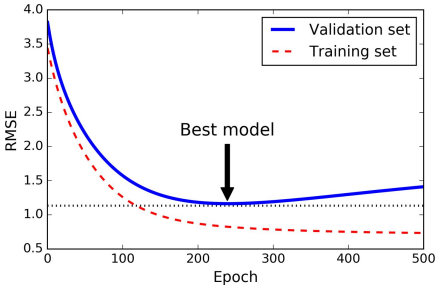

Reference:

1. Hands-On Machine Learning with Scikit-Learn and TensorFlow by Aurélien Géron
2. Why 50% when using dropout?

Phase I - Underfit

Input layer	(n_features, n_prev)
Hidden layer	LSTM
Hidden layer	LSTM (# of neuron = 300)
Output layer	2 (London) 3 (Beijing)
SMAPE (on original data)	0.650053029710999 46

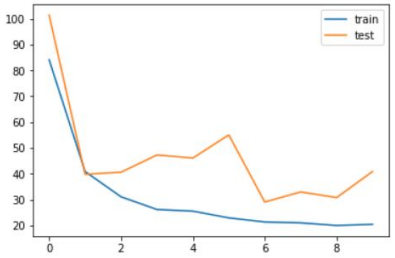

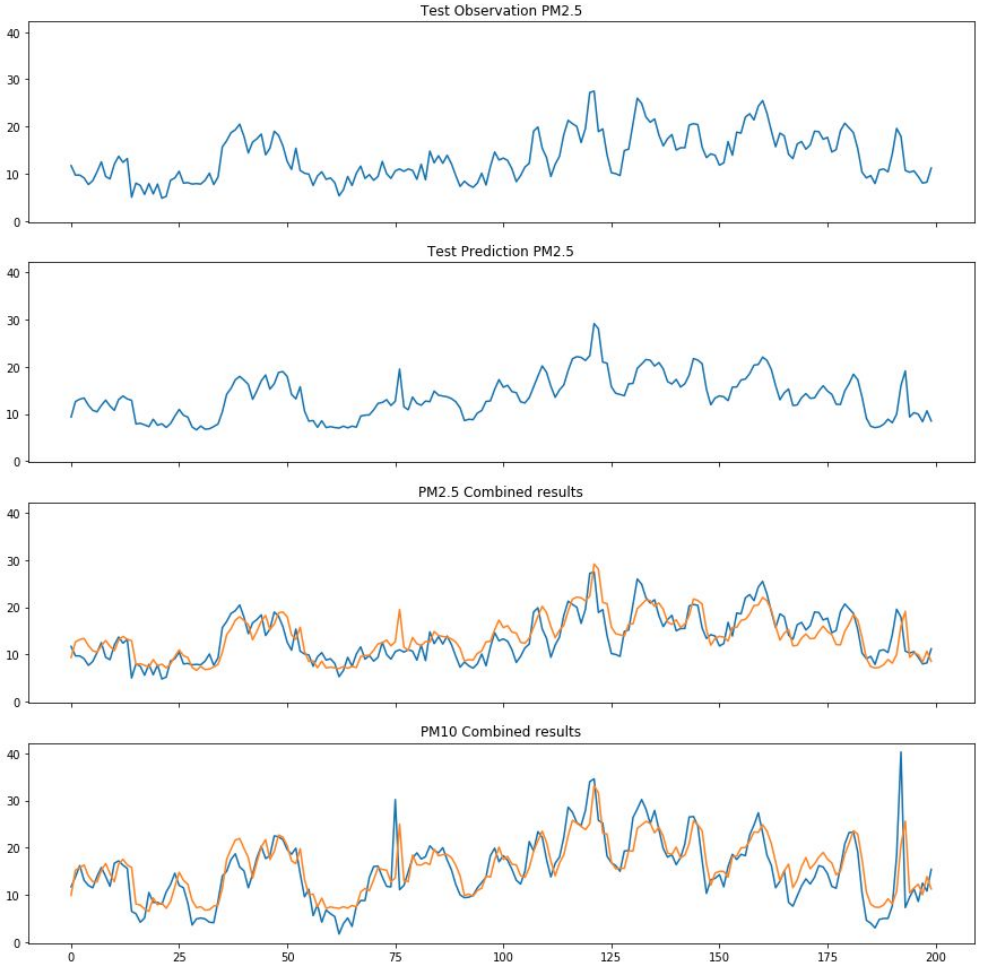

Model Design

- LSTM (GRU is not so good in this competition…)
- Seq2seq

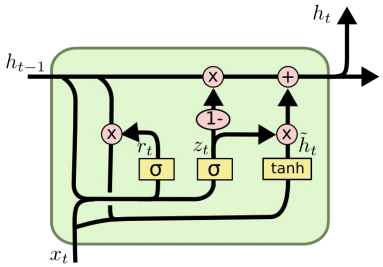

$$z_t = \sigma\left(W_z \cdot [h_{t-1}, x_t]\right)$$
$$r_t = \sigma\left(W_r \cdot [h_{t-1}, x_t]\right)$$
$$\tilde{h}_t = \tanh\left(W \cdot [r_t * h_{t-1}, x_t]\right)$$
$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

LSTM

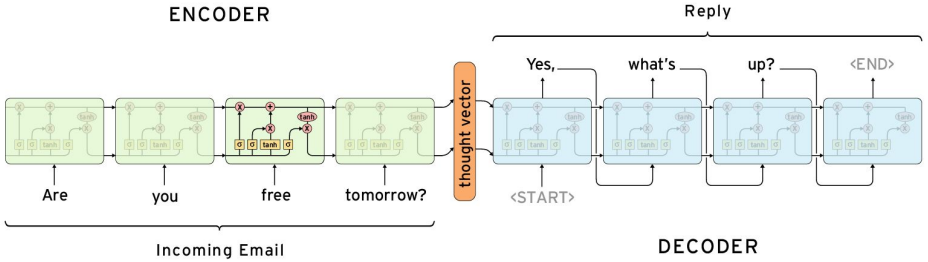

Seq2Seq Model

Reference
[1] Understanding LSTM Networks
[2] Generative Model Chatbots
[3] Sutskever, Ilya, Oriol Vinyals, and Quoc V. Le. "Sequence to sequence learning with neural networks." Advances in neural information processing systems. 2014.
[4] Zaytar, Mohamed Akram, and C. E. El Amrani. "Sequence to sequence weather forecasting with long short term memory recurrent neural networks." Int J Comput Appl 143.11 (2016).

Model Design (cont.)

- Model: Seq2Seq Model
- Layers: 4
- Number of LSTM neurons: 1000

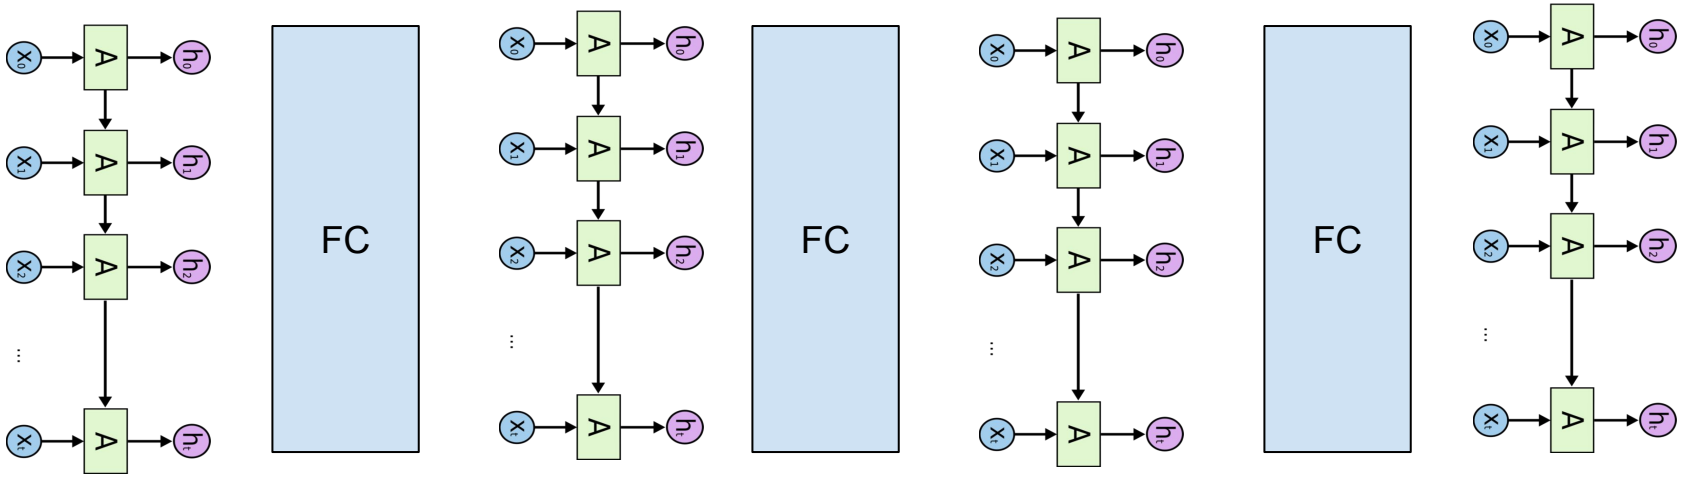

Model Design (cont.) - Bidirectional RNN

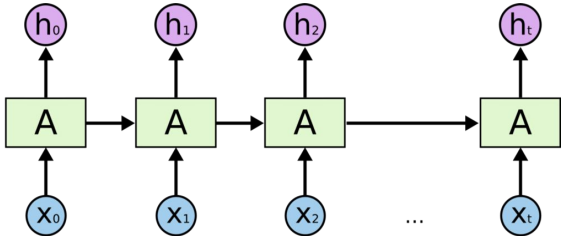

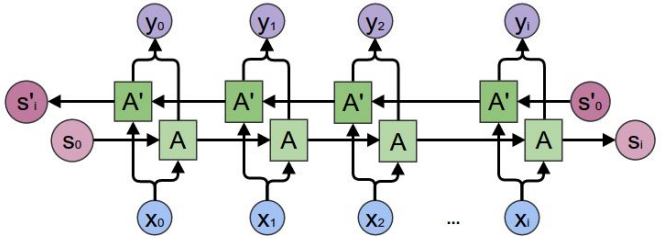

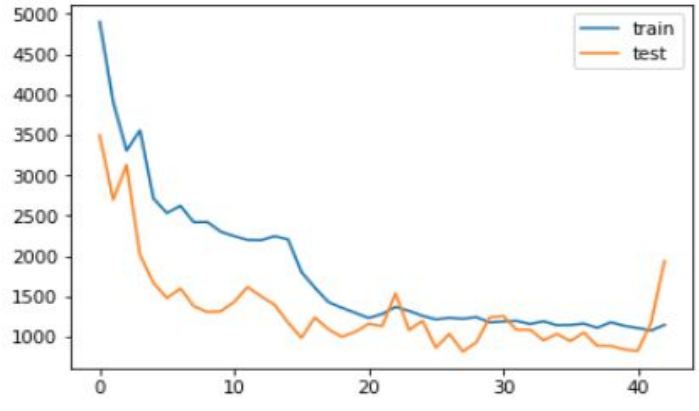

SMAPE: 0.5227025467578719

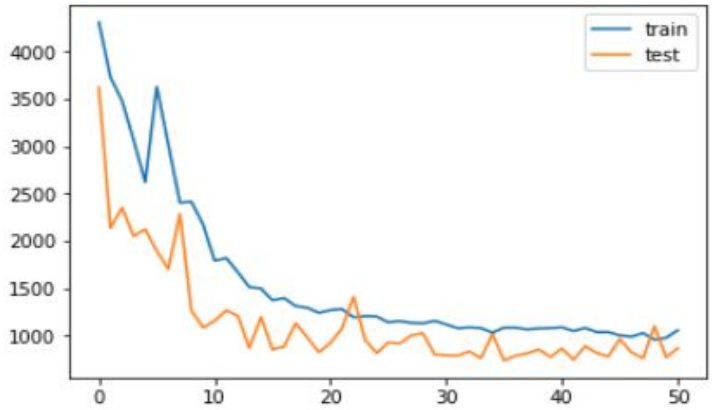

SMAPE: 0.4267148125612628

Model Design (cont.) - Optimizer

Q: Batch Normalization Layer?

Ans: On RNN is a disaster, because the RNN often vary with the length of input sequence. SMAPE got a big escillation in test error

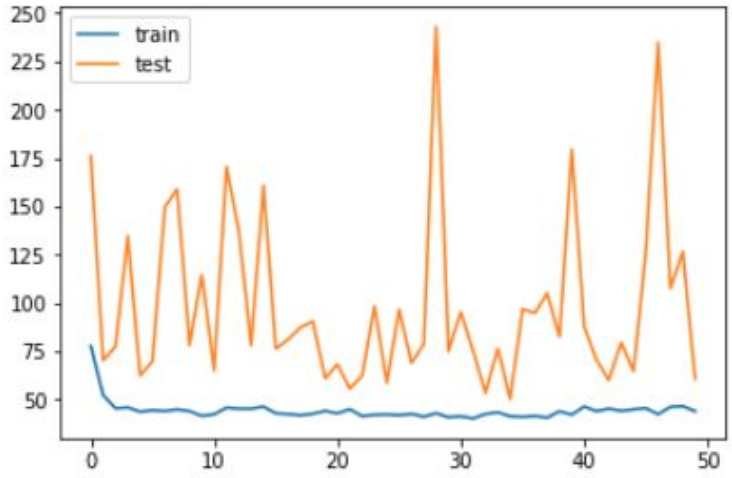

Reference:

1. Quora: Why is it difficult to apply batch-normalization to RNNs?, https://www.quora.com/Why-is-it-difficult-to-apply-batch-normalization-to-RNNs
2. Is it normal to use batch normalization in RNN/lstm RNN?, https://stackoverflow.com/questions/45493384/is-it-normal-to-use-batch-normalization-in-rnn-lstm-rnn

Model Design (cont.)

- Optimizer - Baseline - ADAM
- How’s other Optimizer (ex. rmsprop, adagrad etc)?

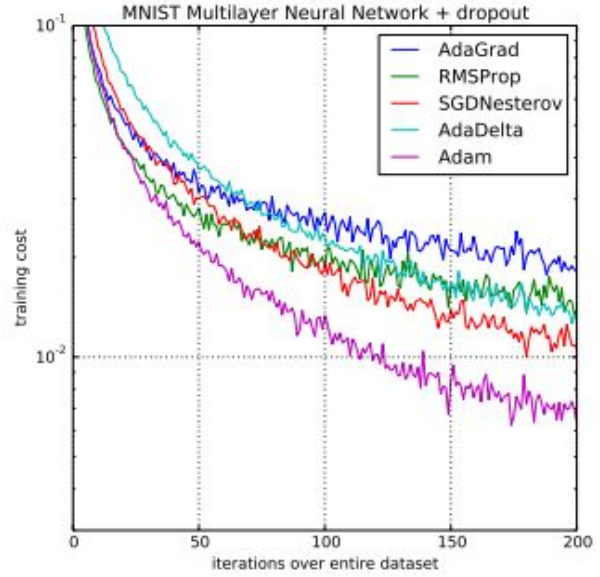

Final Model

Layer (type)	Output Shape	Param #
bidirectional_1 (Bidirection	(None, 100, 2000)	8080000
dense_1 (Dense)	(None, 100, 1000)	2001000
activation_1 (Activation)	(None, 100, 1000)	0
bidirectional_2 (Bidirection	(None, 100, 2000)	16008000
dense_2 (Dense)	(None, 100, 1000)	2001000
activation_2 (Activation)	(None, 100, 1000)	0
dropout_1 (Dropout)	(None, 100, 1000)	0
bidirectional_3 (Bidirection	(None, 100, 2000)	16008000
dense_3 (Dense)	(None, 100, 1000)	2001000
activation_3 (Activation)	(None, 100, 1000)	0
bidirectional_4 (Bidirection	(None, 2000)	16008000
dense_4 (Dense)	(None, 1000)	2001000
activation_4 (Activation)	(None, 1000)	0
dropout_2 (Dropout)	(None, 1000)	0
dense_5 (Dense)	(None, 1000)	1001000
dense_6 (Dense)	(None, 1000)	1001000
dense_7 (Dense)	(None, 3)	3003

Total params: 66,113,003
Trainable params: 66,113,003
Non-trainable params: 0

Reproducible Results - Random seed

- Initialization random weights of Neural Net Model
- Given same initial to keep same random value (start from same position)

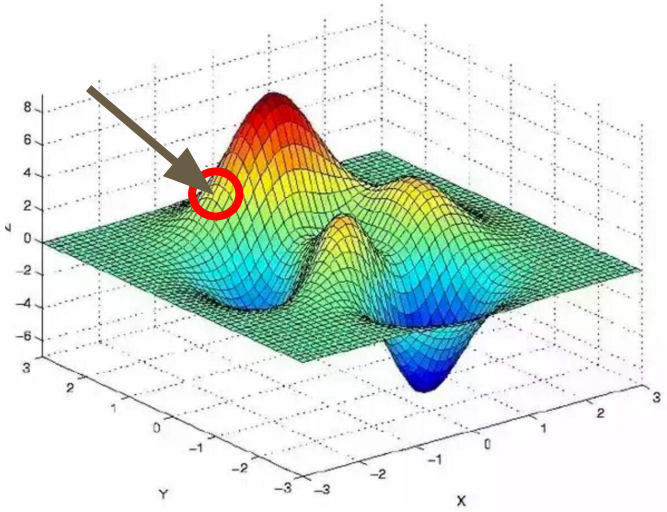

Reference:

1. How to Get Reproducible Results with Keras
2. Why Do I Sometimes Get Better Accuracy With A Higher Learning Rate Gradient Descent?

Results

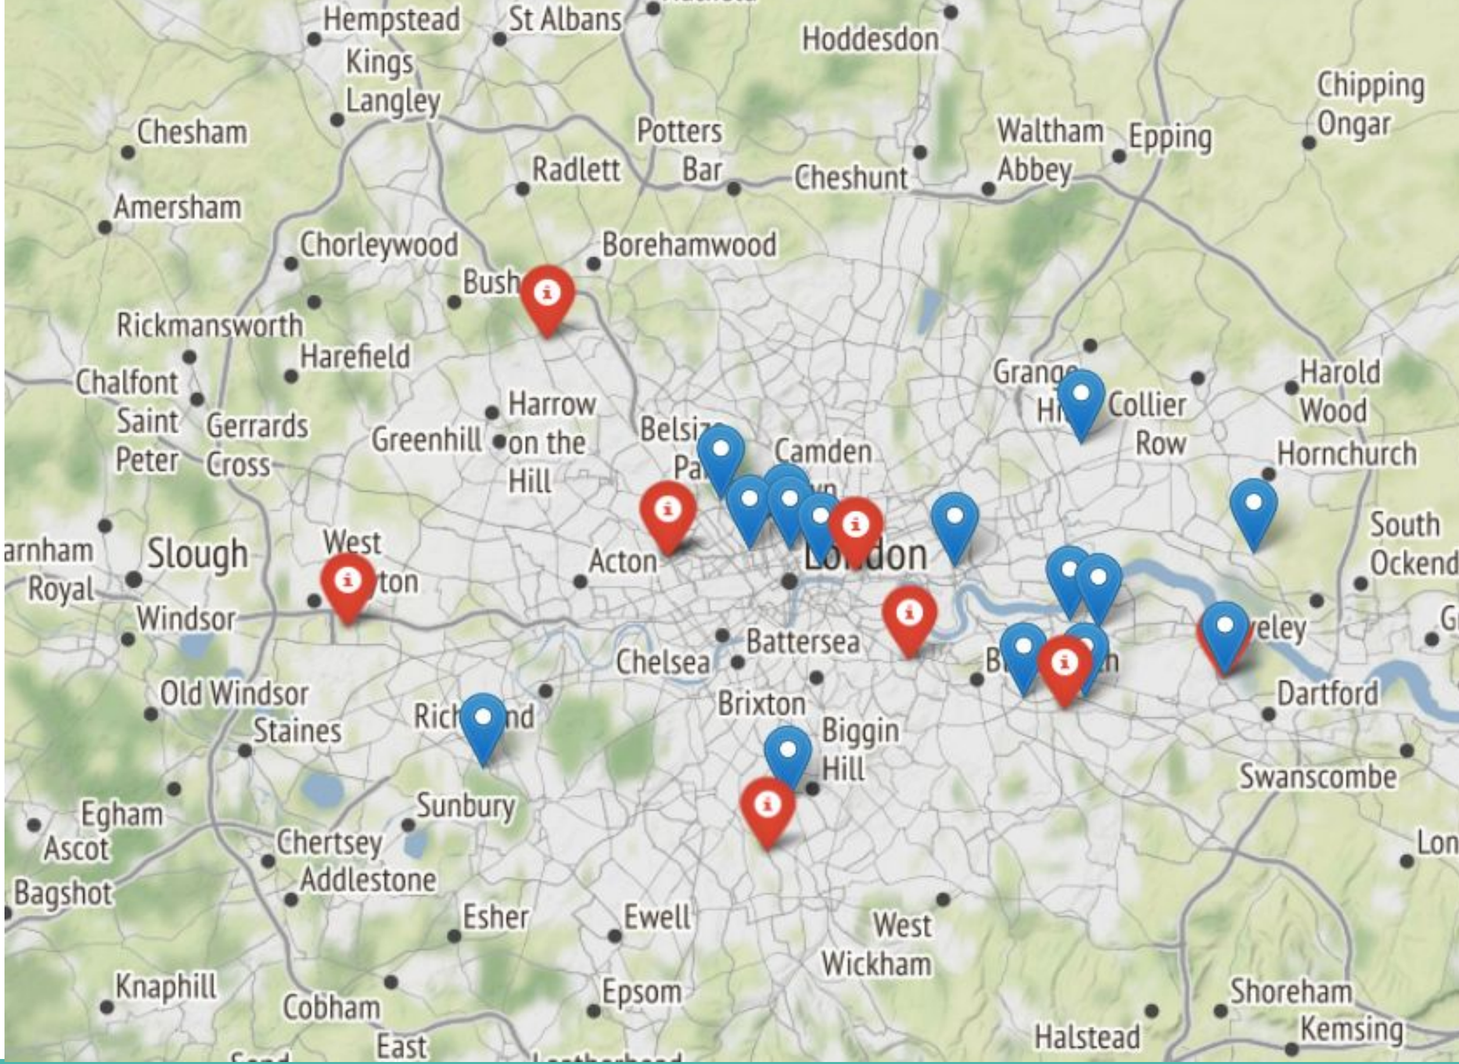
Hempstead
St Albans
Hoddesdon
Kings
Langley
Chipping
Ongar
Chesham
Potters
Bar
Waltham
Abbey
Epping
Radlett
Cheshunt
Amersham
Chorleywood
Borehamwood
Bush
Rickmansworth
Harefield
Chalfont
Saint
Peter
Gerrards
Cross
Harrow
on the
Hill
Greenhill
Camden
Collier
Row
Harold
Wood
Hornchurch
Slough
West
Acton
Royal
Windsor
South
Ockend
Battersea
Chelsea
Brixton
Old Windsor
Staines
Dartford
Biggin
Hill
Sunbury
Swanscombe
Egham
Ascot
Chertsey
Addlestone
Bagshot
Esher
Ewell
West
Wickham
Knaphill
Epsom
Shoreham
Cobham
Kemsing
Halstead
East

Training Time

- 48 Models for every station
- Batch size: 32
- Average Time
  - 1 hr ~ 1.5 hr each city
- Total Time
  - around 2 ~ 3 days

Analysis

- RNN Model Underfit Solution
  - Add more layer
  - Add more hidden unit
  - Fully Connected Layer stabilized model when hidden neuron is a large number
    - SMAPE drop from 1.X -> 0.5 ~ 0.8

Analysis (cont.)

- SMAPE baseline is around 0.75 ~ 0.8
- SMAPE of 2 cities:
  - **Beijing**: [0.3, **0.85**]
  - London: [0.4, 0.57]

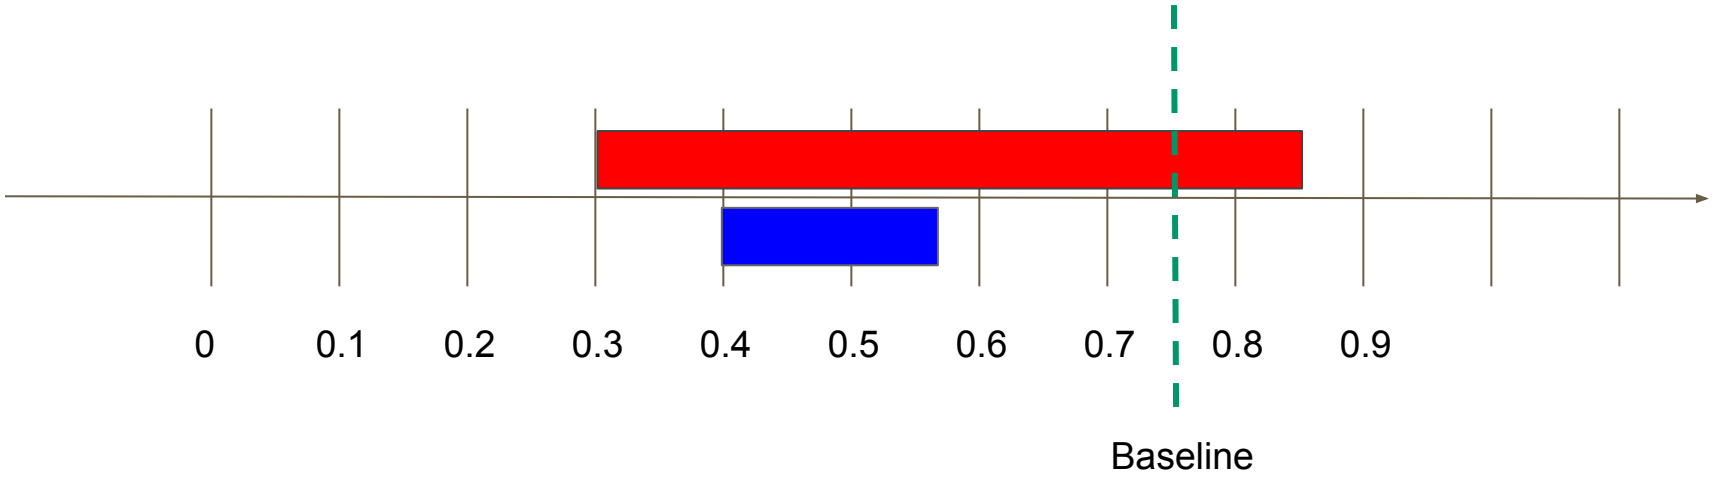

Analysis (cont.)

- PM 2.5 - Official data statistics
  - “...北京市全年PM2.5主要来源中本地排放占三分之二，区域 传输占三分之一[1]”
  - Car emissions gas: CO, NO2 [2]

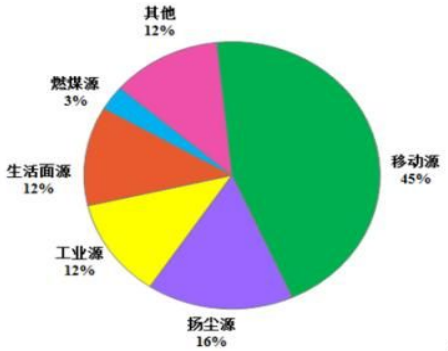

图 2 现阶段北京市大气 $PM_{2.5}$ 本地来源

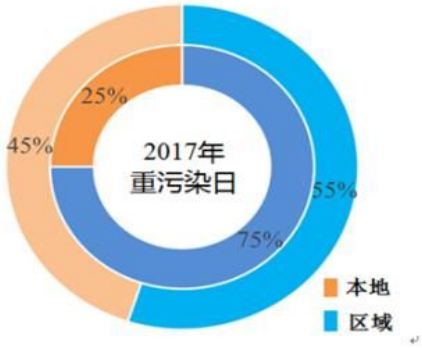

图 1 2017 年北京市重污染日大气 $PM_{2.5}$ 本地和区域贡献

Reference:

1. 最新科研成果新一轮北京市PM2.5来源解析正式发布, Beijing Environmental Protection Bureau
2. Wiki -- Exhaust gas

Analysis (cont.)

- Correlation: feature importance analysis
- Batch size: 64
- Drop SO2 -> SMAPE no too much difference
  - Save time

Correlation	Feature 1	Feature 2	
Positive (+)	PM2.5	CO	around 0.75
	PM2.5	PM2.5	around 0.5
X	PM10	O3	
Negative (-)	O3	NO2	around - 0.5

Beijing city feature correlation plot

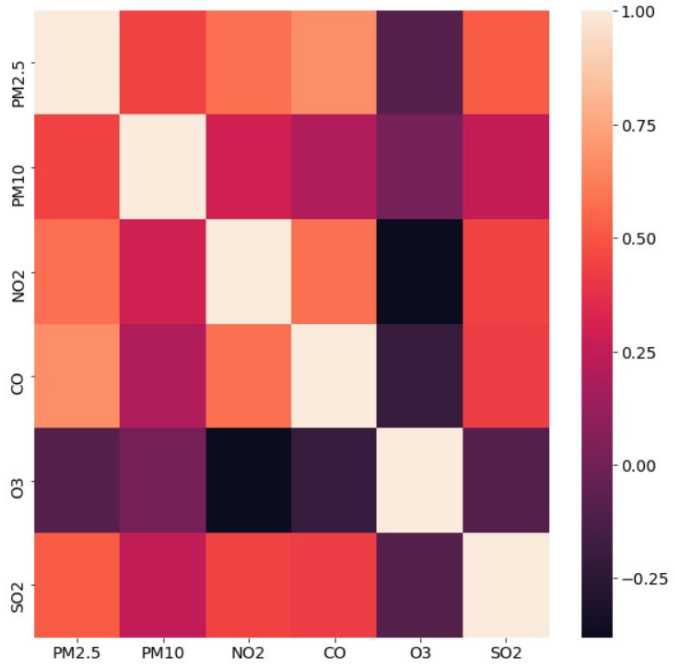

Analysis (cont.)

- Batch Size
  - 32
  - How about larger? 64 or 128?
- SMAPE issue
  - batch size (**up**), smape (**down**)
  - Beijing: [0.45, **1.X**]
  - London: [0.65, 0.8]

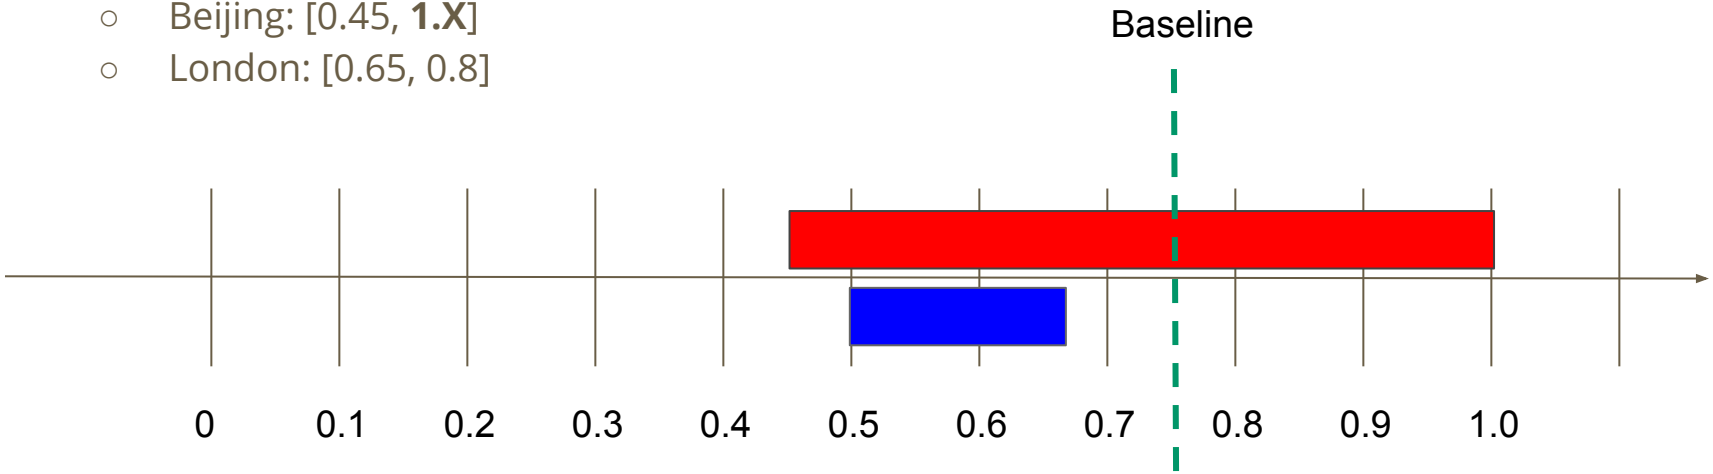

Analysis (cont.)

Classification

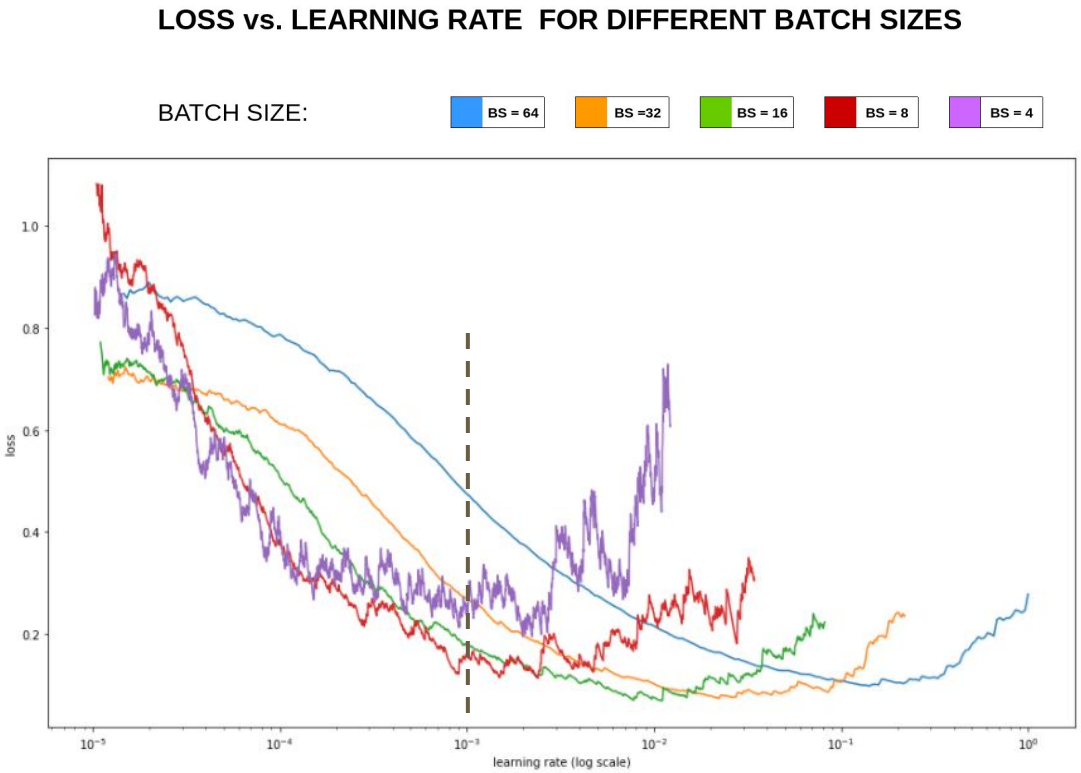

Reference

[1] Visualizing Learning rate vs Batch size

Discussion

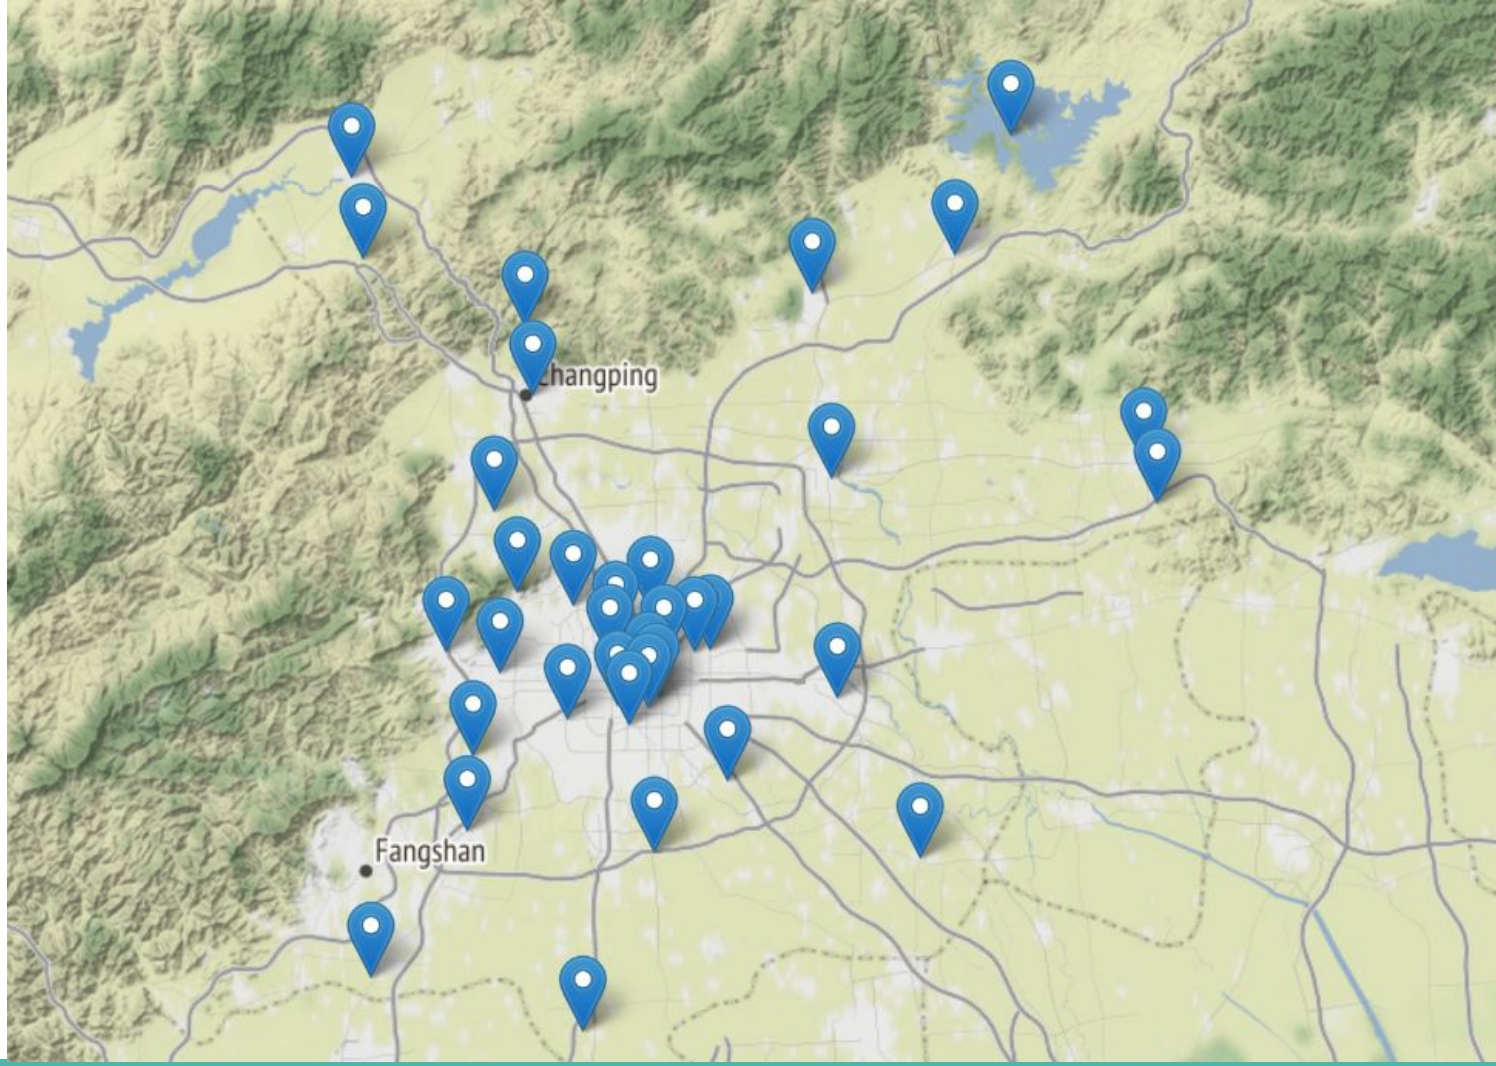

Discussion

- Food(Data) for big guy(RNN) may be not enough
  - Seems the data in this competion is a little too few for the models like RNN. [1]
  - Web Traffic Time Series Forecasting - data set range almost **10** years

Reference

[1] Seq2Seq with attention model and not good score, 92th solution

Lesson Learned...

About competition

- Do this as early as possible
- Data Preprocessing & Cleansing consumes up to **60 - 70%** of time
- Fancy is not always good (A non-convergent Model is a nightmare)

Coworking

- One man do it all ~

Future Plan

- More team member
- Data Preprocessing

Other Resource

1. Kaggle Competition - Web Traffic Time Series Forecasting
2. Arturus/kaggle-web-traffic

Reference

- Géron, Aurélien. Hands-on machine learning with Scikit-Learn and TensorFlow: concepts, tools, and techniques to build intelligent systems. " O'Reilly Media, Inc.", 2017.
- VanderPlas, Jake. Python data science handbook: Essential tools for working with data. " O'Reilly Media, Inc.", 2016.
- Müller, Andreas C., and Sarah Guido. Introduction to machine learning with Python: a guide for data scientists. " O'Reilly Media, Inc.", 2016.
- Sutskever, Ilya, Oriol Vinyals, and Quoc V. Le. "Sequence to sequence learning with neural networks." Advances in neural information processing systems. 2014.
- Zaytar, Mohamed Akram, and C. E. El Amrani. "Sequence to sequence weather forecasting with long short term memory recurrent neural networks." Int J Comput Appl 143.11 (2016).

Thank You For Your Attention

Q&A